

2.uniformed search ("uniformed " means only have successors() function But not which non-goal states )

- **Problem solving as search** : A search problem has five components: S, I, G, actions, cost)

The process to look for sequence of actions to reach the goal

- **problem representation in terms of states**

- Fully observable assumption: no missing information
- -- State Space (aka Problem Space) = all possible valid configurations of the environment

- **state-space graph search formulation: (Formalizing Search in a State Space)**

- Frontier = {S}, where S is the start node
- Loop do
  - if Frontier is empty then return failure
  - pick a node, n, from Frontier
  - if n is a goal node then return solution
  - Generate all n's successor nodes and add them all to Frontier
  - Remove n from Frontier

- **closed world assumption:**

- All necessary information about a problem domain is accessible so that each state is a complete description of the world;

- **expanding a node:** a)generate successor b) add them and their arcs to the state space search tree

- **frontier/open list:** The generated, but not yet expanded, states define the Frontier (aka Open or Fringe) set

- Different Search Strategies: The essential difference is, which state in the Frontier to expand next)

- **partial solution path:** a partial solution path from the start node to the given nodee

- **chronological backtracking:** when search hits a dead end, backs up one level at a time when search hits a dead end,

– problematic if the mistake occurs because of a bad action choice near the top of search tree

- **detecting repeated states,(If State Space is Not a Tree):**

When pick a node from Frontier – Remove it from Frontier – Add it to Explored – Expand node, generating all successors

4.local search: operate using a single **current node** and generally move only to neighbors of that node. greedy local search every node is a solution

- **escaping local optimal**

Fix by replacing fixed probability, p, that a bad move is accepted, with a probability that decreases as the search proceeds

- **Boltzman's equation**

- Let  $\Delta E = f(\text{newNode}) - f(\text{currentNode}) < 0$        $p = e^{-(\Delta E / T)}$  (Boltzman's equation\*)

- **cooling schedule**

- --T, the annealing "temperature," is the parameter that control the probability of taking of bad steps
- --We gradually reduce the temperature, T(k)
- --At each temperature, the search is allowed to proceed for a certain number of steps, L(k)
- --The choice of parameters {T(k),L(k)} is called the cooling schedule

Complete : a complete algorithm will find a solution)

Optimal/admissible : an admissible algorithm will find a solution with minimum cost)

Time complexity: measured for worst case)

Space complexity: maximum size of frontier)

## 7.supervised learning

- **K-nearest neighbor algorithm,**

- **Decision Trees**

- **Ockham's razor, (the Preference Bias of best decision tree is Ockham's razor:)**

- -- The simplest hypothesis that is consistent with all observations is most likely
- -- The smallest decision tree that correctly classifies all of the training examples is best

- **decision tree algorithm,**

- **information gain,** (Goal: Select the attribute that will result in the smallest expected tree size)

- $I(Y; X) = H(Y) - H(Y | X)$

- **max-gain,:** the attribute that has the largest expected information gain

- **entropy,**

- **conditional entropy,**

- **remainder,**  $H(Y|X) = \text{Remainder}(X)$

- **overfitting problem:** As d increases, the(Mean Squared Error ) on the training data improves, but prediction on test data worsens

- – meaningless regularity is found in the data
- – irrelevant attributes confound the true, important, distinguishing features
- – fix by pruning some nodes in the decision tree

- **Pruning**

- Randomly split the training data into TRAIN and TUNE, say 70% and 30% 2.
- Build a full tree using only the TRAIN set 3.
- Prune the tree using the TUNE set,

- **setting parameters,**

- – most learning algorithms require setting various parameters – they must be set without looking at the Test data
- – Common approach: use a Tuning Set
- 1. Partition given examples into TRAIN, TUNE and TEST sets
- 2. For each candidate parameter value, generate a decision tree using the TRAIN set
- 3. Use the TUNE set to evaluate error rates and determine which parameter value is best
- 4. Compute the final decision tree using the selected parameter values and both TRAIN and TUNE sets
- 5. Use TEST to compute performance accuracy

3.informed search(use domain knowledge to guide selection of the best path to continue searching)

- **Heuristic function:** estimated cost of the cheapest path from the state at node n to a goal state. h(n)

- Evaluation function: cost estimate,

- **admissible heuristic:** *never overestimates* the cost to reach the goal.

- **consistent heuristic,**  $h(n) \leq c(n, n') + h(n')$

- **devising heuristics: relaxing the problem; If optimality is not required, a satisficing solution is okay, (fewer expand nodes)**

Goal of the heuristic is then to get as close as possible, either under or over, to the actual cost)

- Heuristics are often defined by relaxing the problem,

- If  $h(n) = h^*(n)$  for all n,

- – only nodes on optimal solution path are expanded
- – no unnecessary work is performed

- If  $h(n) = 0$  for all n,

- – the heuristic is admissible
- – A\* performs exactly as Uniform-Cost Search (UCS)

- If  $h_1(n) \leq h_2(n) \leq h^*(n)$  for all n, then h2 dominates h1

- – h2 is a better heuristic than h1

- – A\* using h1 (i.e., A1\*) expands at least as many if not more nodes than using A\* with h2

## 5. game playing (should consider opponent)

- **Zero-sum games:** one player's gain is the other player's loss. Does not mean fair

- **perfect information games:** each player can see the complete game state. No simultaneous decisions

- **deterministic vs. stochastic games**

**deterministic;** fully observable environments where two agents act alternately and at the end of the game are always equal and opposite

**Stochastic:** including a random element

- **game playing as search** : . (represent both computer and opponent moves)

Initial state; player; action; result; terminal state;

- **ply:** in minimal , consisting of two half-moves, each of which is called a **ply**

- **minimax principle,(cannot be used to end of the game)**

- high values favor the computer -- The computer assumes after it moves the opponent will choose the minimizing move

- low values favor the opponent -- The computer chooses the best move considering both its move and the opponent's optimal move

- **static evaluation function:** use heuristics to estimate the value of non-terminal states (agree with the Utility function when calculated at terminal nodes)

- **alphabeta pruning,** ( Pruning can be used to ignore some branches)

- **iterative-deepening with alpha-beta (for Dealing with Limited Time- > time can also be saved by limited depth)**

- – run alpha-beta search with depth-first search and an increasing depth-limit
- – when time runs out, use the solution found for the last completed alpha-beta search
- – "anytime algorithm"

- **horizon effect** :The computer has a limited horizon, it cannot see that this significant event could happen

- **quiescence search —> used for avoid "short sightedness"(can also be solved by secondary search)**

- – when SBE value is frequently changing, look deeper than the depth-limit
- – look for point when game "quiets down"
- – E.g., always expand any forced sequences

- **representing non-deterministic games( games involve chance,)**

- **chance nodes in expectimax algorithm:** representing random events

## 6.unsupervised learning

- **Inductive learning problem,:** Generalize from a given set of (training) examples, make predictions for future

- **unsupervised learning problem,** :the agent learns patterns in the input even though no explicit feedback is supplied

- **Feature:** Each dimension in x is called a **feature or attribute**. --- x is a point in the D-dimensional **feature space**

- **Labels:** the desired prediction for an instance x

- **classes,:** Discrete labels:

- **classification problems,:** if y discrete

- **inductive bias, :**

- Inductive Bias – is used when one h is chosen over another – is needed to generalize beyond the specific training examples

- **Completely unbiased inductive algorithm** – only memorizes training examples – can't predict anything about unseen examples

- **preference bias:** define a metric for comparing h's so as to determine whether one is better than another

- **hierarchical agglomerative clustering algorithm,** : Build a binary tree over the dataset by repeatedly merging clusters

- **single linkage, / complete linkage, /average linkage,**

- shortest distance /largest distance/average distance between all pairs of members, one from each cluster

- **dendrogram,:** The binary tree get in **hierarchical agglomerative clustering** is often called a **dendrogram**, or taxonomy, or a hierarchy of data points

- **kmeans clustering algorithm,** :Specify the desired number of clusters and use an iterative algorithm to find them

Stop until cluster centers no longer change. Will terminate! Not optimal

- **cluster center,:**Choose ci to be the mean / centroid of all points/examples in the cluster

- **distortion cluster quality.** (Distortion = Sum of squared distances of each data point to its cluster center)

	description	Compl ete	Optim al/	Time complexity		Space complexit	nodes	
BFS		Criterion	Breadth- First	Uniform- Cost	Depth- First	Depth- Limited	Iterative Deepening	Bidirectional (if applicable)
DFS		Complete?	Yes <sup>a</sup>	Yes <sup>a,b</sup>	No	No	Yes <sup>a</sup>	Yes <sup>a,d</sup>
		Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$	$O(b^{d/2})$
		Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$	$O(b^{d/2})$
		Optimal?	Yes <sup>c</sup>	Yes	No	No	Yes <sup>c</sup>	Yes <sup>c,d</sup>
<b>Figure 3.21</b> Evaluation of tree-search strategies. <i>b</i> is the branching factor; <i>d</i> is the depth of the shallowest solution; <i>m</i> is the maximum depth of the search tree; <i>l</i> is the depth limit. Superscript caveats are as follows: <sup>a</sup> complete if <i>b</i> is finite; <sup>b</sup> complete if step costs $\geq \epsilon$ for positive $\epsilon$ ; <sup>c</sup> optimal if step costs are all identical; <sup>d</sup> if both directions use breadth-first search.								
IDS	do DFS to depth 1 and treat all children of the start node as leaves – if no solution found, do DFS to depth 2							
bidirectional search	bfs from both start and goal --Stop when Frontiers meet						$O(b^{(d/2)})$	
Greedy bfs	Use as an evaluation function, $f(n) = h(n)$ , sorting nodes in the Frontier	no	no				Deterministic	
Beam search	Use an evaluation function $f(n) = h(n)$ as in Greedy bfs, and restrict the maximum size of the Frontier to a constant, k	no	no				Fully observable	Checkers, Chess, Go, Othello
							Partially observable	stratego, battleship
Best-first search	Sort nodes in the Frontier list by increasing values of an evaluation function	No	No				$b^m$	$b^m$
Algorithm a	Use as an evaluation function $f(n) = g(n) + h(n)$ ,	no						
A *	all nodes n in the search space, $h(n) \leq h^*(n)$ should terminate only when a goal is removed from the priority queue	yes	yes				$O(\text{NO. states})$	
Hill climbing	1.Pick initial state s 2. Pick t in neighbors(s) with the largest f(t) 3. if $f(t) \leq f(s)$ then stop and return s							
hill-climbing with random restarts	When stuck, pick a random new starting state and re-run hill-climbing from there							
stochastic hill-climbing (less greedy)	1.Pick initial state, s 2. Randomly pick state t from neighbors of s 3. if f(t) better than f(s) then $s = t$ // t better than s so move there else with small probability $s = t$ 4. Goto Step 2 until some stopping criterion is met							
minimax algorithm					$O(b^d)$		$O(bd)$	
alpha-beta pruning	Worst Case: – ordered so that no pruning takes place Best Case: – each player's best move is visited first				In practice often get $O(b^{(d/2)})$			
Expectimax	For chance nodes: compute the average, weighted by the probabilities of each child							

Exploitation term

$$\frac{w_i}{n_i} + c \sqrt{\frac{\ln t}{n_i}}$$

Exploration term

**Hill climbing:** Solution found by HC is totally determined by the starting point  
 When the neighbor is large: • Randomly generate neighbors, one at a time • If neighbor is better, take the move  
**Simulated Annealing:** fast, increase chance to find global optimum  
 As  $\Delta E \rightarrow -\infty$ ,  $p \rightarrow 0$  i.e., as move gets worse, probability of taking it decreases exponentially •  
 As  $T \rightarrow 0$ ,  $p \rightarrow 0$  i.e., as “temperature,” T, decreases probability of taking bad move decreases  
 If  $\Delta E \ll T$  if badness of move is small compared to T, move is likely to be accepted  
 If  $\Delta E \gg T$  if badness of move is large compared to T, move is unlikely to be accepted

- Tips:
- starting temp must be high to escape local optimal
  - With an infinitely slow cooling rate, SA finds the global optimum with probability 1

**Effectiveness of Alpha-Beta Search:** depends on the order in which successors are examined  
 Worst Case: ordered so that no pruning takes place  
 Best Case: – each player's best move is visited first

**Non-Deterministic Games:** increases branching factor;  
 Value of look-ahead diminishes: less effective for alpha-beta pruning  
 Improve performance in game playing: reduce depth (better search, learn good features); reduce breadth (explore subset of possible moves use randomized exploration)

**Monte Carlo tree search** (Best-first search based on random sampling of search space; focus on good moves)  
 simulate k random games by selecting moves at random for both players until game over (called playouts); count how many were wins out of each k playouts; move with most wins is selected (infinite moves and infinite length)  
 Recursively build search tree, where each round consists of:

- Selection: Starting at root, successively select best child nodes using scoring method until leaf node L reached
- Expansion: Create and add best (or random) new child node, C, of L
- Simulation: Perform a (random) playout from C
- Backpropagation: Update score at C and all of C's ancestors in search tree based on playout results

**Smarter Initialization of K-Means Clusters:**  
 1. Run k-Means multiple times with different starting, random cluster centers  
 2. \*\*pick a random point x1 from the dataset 1. Find a point x2 far from x1 in the dataset  
**Knn:** pick k  
 Split data into training and tuning sets; Classify tuning set with different values of k; Pick the k that produces the smallest tuning-set error

**K fold cross evaluation:**  
 1. Divide all examples into K disjoint subsets  $E = E_1, E_2, \dots, E_K$   
 2. For each  $i = 1, \dots, K$ : let TEST set =  $E_i$  and TRAIN set =  $E - E_i$ ; build decision tree using TRAIN set; determine accuracy Acc<sub>i</sub> using TEST set  
 3. Compute K-fold cross-validation estimate of performance = mean accuracy  
**leave\_one\_out cross evaluation:** (use when have small dataset)  
 For  $i = 1$  to N do // N = number of examples  
 1. Let  $(x_i, y_i)$  be the i th example  
 2. Remove  $(x_i, y_i)$  from the dataset  
 3. Train on the remaining N-1 examples  
 4. Compute accuracy on i th example (Accuracy = mean accuracy on all N runs)  
**Random forest:** collection of independently-trained binary decision trees  
 Hard to interpret compared to decision tree, but efficient and can handle large data  
 1. For each tree, pick a randomly sampled subset of training data  
 2. Randomly choose a subset of features and thresholds at each node  
 3. Pick the feature and threshold that give the largest information gain  
 4. Recurse until a certain accuracy is achieved or depth-bound reached

**Ensemble strategies**

- Boosting
- Bagging

Given N training examples, generate separate training sets by choosing n examples with replacement from all N examples • Called “taking a bootstrap sample” or “randomizing the training set” • Construct a separate classifier using the n examples in each training set

**BFS:** if all operators (i.e., arcs) have the same constant cost, or costs are positive, non-decreasing with depth – otherwise, not optimal but does guarantee finding solution of shortest length (i.e., fewest arcs)

**DFS:** Can find long solutions quickly if lucky) (May not terminate without a depth bound)

**IDS:** optimality as stated for BFS . It Trades a little time for a huge reduction in space; “Anytime” algorithm: good for response-time critical applications

Has advantages of BFS – completeness – optimality as stated for BFS

Has advantages of DFS – limited space – in practice, even with redundant effort it still finds longer paths more quickly than BFS

**stochastic hill-climbing:** bad move: when probability that decreases as the search progress (T decreases) and with the badness of move ( $\Delta E$  worsens)

**A\*:** the tree-search version of A\* is optimal if  $h(n)$  is admissible, while the graph-search version is optimal if  $h(n)$  is consistent. (big search space  $\rightarrow$  iterative deepening A\*)

**Simulated annealing:** The probability decreases : 1) with the “badness” of the move—the amount  $\Delta E$  by which the evaluation is worsened. 2) “temperature” T goes down: “bad” moves are more likely to be allowed at the start when T is high, and they become more unlikely as T decreases. If the schedule lowers T slowly enough, the algorithm will find a global optimum with probability approaching 1